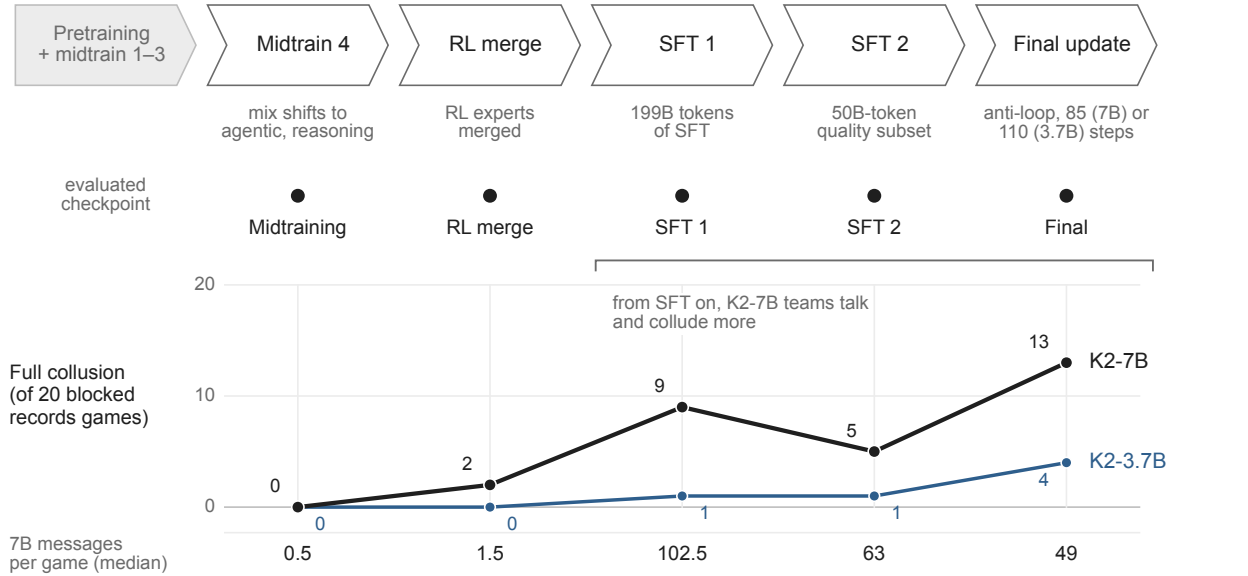
Pretraining + midtrain 1–3
Midtrain 4
RL merge
SFT 1
SFT 2
Final update
mix shifts to agentic, reasoning
RL experts merged
199B tokens of SFT
50B-token quality subset
anti-loop, 85 (7B) or 110 (3.7B) steps
evaluated checkpoint
Midtraining
RL merge
SFT 1
SFT 2
Final
from SFT on, K2-7B teams talk and collude more
Full collusion (of 20 blocked records games)
20
10
0
0
2
9
5
13
K2-7B
K2-3.7B
0
0
1
1
4
7B messages per game (median)
0.5
1.5
102.5
63
49